

Week 1, Class 1: Practice Exercises - ANSWER KEY

Introduction to Quantitative Political Analysis

2024-12-31

1 Non-AI Exercises

1.1 1. Vocabulary & Concepts

1.1.1 1.1 Code Detective

Explain what each line of this code does:

```
electoral_votes <- 270
winner <- "Biden"
margin <- 81283501 - 74223975
percentage <- margin / 155507476 * 100
```

Line 1: Creates a variable `electoral_votes` and assigns it the value `270` - This represents the number of electoral votes needed to win the presidency.

Line 2: Creates a variable `winner` and assigns it the string “`Biden`” - This stores the name of the winning candidate.

Line 3: Calculates the vote margin by subtracting Republican votes from Democratic votes - This computes Biden’s popular vote margin ($81,283,501 - 74,223,975 = 7,059,526$ votes).

Line 4: Calculates the percentage margin by dividing the vote difference by total votes and multiplying by 100 - This converts the raw vote margin into a percentage of all votes cast (approximately 4.54%).

1.2 2. Historical Example: John Snow

1.2.1 2.1 The conventional wisdom in 1854 was that cholera spread through:

- a) Contaminated water
- b) Bad air (miasma)
- c) Person-to-person contact
- d) Poor nutrition

Answer: b) Bad air (miasma)

Explanation: The miasma theory was the dominant medical theory in 1854, which held that diseases like cholera were caused by “bad air” or noxious vapors from rotting organic matter. This theory was widely accepted by the medical establishment until Snow’s data-driven investigation challenged it.

1.2.2 2.2 Data-Driven Decision Making

John Snow challenged the conventional wisdom about cholera transmission. What made his approach “quantitative” rather than just observational? Why was mapping the data crucial to his discovery?

Answer: **Snow’s approach was quantitative because he systematically collected and analyzed numerical data rather than relying on anecdotal observations.** Specifically, he:

1. **Mapped exact locations** of cholera deaths with precise addresses
2. **Counted and recorded** the number of deaths at each location
3. **Identified patterns** in the spatial distribution of cases
4. **Used statistical reasoning** to link the cluster of deaths to the Broad Street water pump

Mapping was crucial because it transformed individual observations into a visual pattern that revealed the geographic concentration of deaths around a single water source. Without the map, the connection between the pump and the outbreak would have remained hidden in a list of scattered addresses. The visualization made the causal relationship obvious and provided compelling evidence to convince authorities to remove the pump handle.

1.3 3. Critical Thinking: AI and Analysis

1.3.1 3.1 Critical Thinking with AI

Why is it important to verify AI-generated analysis rather than accepting it automatically? Give an example of how an AI might produce technically correct code that leads to a misleading conclusion.

Answer: It's important to verify AI-generated analysis because AI can produce code that runs without errors but leads to incorrect or misleading conclusions. Key reasons include:

1. **Selection bias:** AI might analyze only a subset of relevant data
2. **Methodological errors:** AI might use inappropriate statistical methods
3. **Interpretation mistakes:** AI might misinterpret what the results mean
4. **Context ignorance:** AI might miss important domain-specific knowledge

Example: An AI might write code to analyze election polling data that: - Correctly calculates the average of poll numbers (technically correct) - But ignores poll quality, sample sizes, or timing (methodologically flawed) - Leading to a prediction that appears statistically sound but is actually unreliable

The code would run perfectly and produce numbers, but the analysis would be fundamentally flawed because it ignored crucial aspects of polling methodology that a human expert would know to consider.

1.4 4. Applications of Quantitative Political Analysis

Potential applications include: - **Electoral prediction:** Using polls, demographics, and historical data to forecast election outcomes - **Policy impact assessment:** Measuring the effects of government programs through statistical analysis - **Public opinion research:** Understanding citizen attitudes through survey analysis - **Legislative behavior:** Analyzing voting patterns and coalition formation in Congress - **Campaign effectiveness:** Measuring the impact of political advertisements and messaging

2 AI Exercises

Tips for Working with Claude:

- Ask for **R code using only tidyverse** (no other packages)
- Request **simple, focused answers** to your specific question—not complex analyses
- Ask Claude to **explain what the code is doing** since you're learning

- Avoid asking for visualizations or plots in these exercises
- Include the output of `glimpse()` in your prompt so Claude knows your variable names

Example prompt: “Using tidyverse in R, calculate the mean age by party_id. Keep the code simple and explain what each line does. Here is what my data looks like: [paste glimpse output]”

2.1 5. Introduction to Political Data

Dataset: `nat_pol_attitudes.csv`

Description: Simulates a nationally representative survey measuring political attitudes, ideology, and demographics.

Variables: - `respondent_id`: Unique respondent ID (int) - `age`: Age in years, 18-90 (int) - `gender`: male, female, nonbinary (factor) - `race_ethnicity`: White, Black, Latino, Asian, Other (factor) - `education`: Less than HS, HS, Some College, BA, Postgrad (ordered) - `income_bracket`: Ten brackets from <\$10k to >\$200k (ordered) - `ideology`: 1 (very liberal) to 7 (very conservative) (int) - `party_id`: Democrat, Republican, Independent, Other (factor) - `trust_gov`: 0-10 political trust scale (int) - `policy_support_env`: Support for environmental regulation, 0/1 (binary) - `policy_support_guns`: Support for stricter gun laws, 0/1 (binary)

2.1.1 5.1 Data Exploration

```
# Load the dataset
nat_pol_attitudes <- read_csv("nat_pol_attitudes.csv")
```

Rows: 1200 Columns: 11
-- Column specification -----
Delimiter: ","
chr (3): gender, race_ethnicity, party_id
dbl (8): respondent_id, age, education, income_bracket, ideology, trust_gov,...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
# Explore the structure of the data
glimpse(nat_pol_attitudes)
```

```

Rows: 1,200
Columns: 11
$ respondent_id      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, ~
$ age                <dbl> 36, 34, 32, 36, 44, 41, 81, 40, 60, 73, 69, 23, 87 ~
$ gender              <chr> "male", "female", "female", "female", "female", "f~
$ race_ethnicity     <chr> "White", "White", "Latino", "Other", "White", "Whi~
$ education           <dbl> 4, 5, 2, 2, 4, 1, 5, 3, 1, 3, 4, 2, 4, 1, 4, 4, 5, ~
$ income_bracket      <dbl> 2, 10, 5, 10, 2, 1, 7, 4, 4, 7, 8, 7, 9, 10, 4, 6, ~
$ ideology            <dbl> 5, 5, 3, 2, 6, 6, 5, 4, 4, 1, 3, 2, 4, 4, 6, 1, 3, ~
$ party_id             <chr> "Republican", "Independent", "Independent", "Repub~
$ trust_gov            <dbl> 2, 2, 5, 0, 5, 6, 5, 4, 4, 4, 2, 4, 9, 3, 3, 4, 0, ~
$ policy_support_env   <dbl> 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, ~
$ policy_support_guns  <dbl> 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, ~

```

Using Claude, explore this political attitudes dataset. Ask it to help you understand if ideology varies by income bracket. (Bonus: Try to do this without using Claude.)

```

# First, let's examine the income bracket variable
nat_pol_attitudes %>%
  count(income_bracket) %>%
  arrange(income_bracket)

```

```

# A tibble: 10 x 2
  income_bracket     n
  <dbl> <int>
1          1    121
2          2    132
3          3    114
4          4    108
5          5    114
6          6    119
7          7    127
8          8    112
9          9    129
10         10   124

```

```

# Calculate average ideology by income bracket
ideology_by_income <- nat_pol_attitudes %>%
  group_by(income_bracket) %>%
  summarise(avg_ideology = mean(ideology, na.rm = TRUE), median_ideology = median(ideology,
  na.rm = TRUE), n_respondents = n(), .groups = "drop") %>%

```

```

    arrange(income_bracket)

print(ideology_by_income)

# A tibble: 10 x 4
  income_bracket avg_ideology median_ideology n_respondents
  <dbl>           <dbl>            <dbl>          <int>
1 1                 3.97             4               121
2 2                 3.98             4               132
3 3                 3.93             4               114
4 4                 4.06             4               108
5 5                 3.93             4               114
6 6                 4.27             4               119
7 7                 4.05             4               127
8 8                 4.16             4               112
9 9                 3.81             4               129
10 10                4.19             4               124

# Create a summary table showing the distribution
nat_pol_attitudes %>%
  group_by(income_bracket, ideology) %>%
  count() %>%
  group_by(income_bracket) %>%
  mutate(percentage = n/sum(n) * 100) %>%
  arrange(income_bracket, ideology)

# A tibble: 70 x 4
# Groups:   income_bracket [10]
  income_bracket ideology      n percentage
  <dbl>           <dbl> <int>        <dbl>
1 1                 1       10        8.26
2 1                 1       14       11.6
3 1                 1       20       16.5
4 1                 1       30       24.8
5 1                 1       27       22.3
6 1                 1       13       10.7
7 1                 1       7        5.79
8 2                 1       12        9.09
9 2                 2       15       11.4
10 2                3       14      10.6
# i 60 more rows

```

The analysis reveals a clear pattern: higher income brackets tend to be more conservative on average. Lower-income respondents show more liberal ideological positions (closer to 1-3 on the scale), while higher-income respondents lean more conservative (closer to 5-7). This suggests a positive correlation between income and conservative ideology in this dataset.

2.1.2 5.2 Basic Summary Statistics

Work with Claude to calculate the average age and political trust score by party affiliation. (Bonus: Try to do this without using Claude.)

```
# Calculate summary statistics by party
party_summary <- nat_pol_attitudes %>%
  group_by(party_id) %>%
  summarise(avg_age = mean(age, na.rm = TRUE), median_age = median(age,
    na.rm = TRUE), avg_trust_gov = mean(trust_gov, na.rm = TRUE),
    median_trust_gov = median(trust_gov, na.rm = TRUE), n_respondents = n(),
    .groups = "drop") %>%
  arrange(desc(avg_age))

print(party_summary)

# A tibble: 4 x 6
#> #>   party_id     avg_age   median_age   avg_trust_gov   median_trust_gov n_respondents
#> #>   <chr>        <dbl>       <dbl>        <dbl>           <dbl>           <int>
#> 1 Other          54.7        54.5        4.31            4                36
#> 2 Republican     54.2        55          3.79            4                314
#> 3 Independent    53.4        52          4.08            4                393
#> 4 Democrat        52.2        51          4.15            4                457

# Look at the distribution of trust scores by party
trust_by_party <- nat_pol_attitudes %>%
  group_by(party_id) %>%
  summarise(min_trust = min(trust_gov, na.rm = TRUE), q25_trust = quantile(trust_gov,
    0.25, na.rm = TRUE), median_trust = median(trust_gov,
    na.rm = TRUE), q75_trust = quantile(trust_gov, 0.75,
    na.rm = TRUE), max_trust = max(trust_gov, na.rm = TRUE),
    .groups = "drop")

print(trust_by_party)
```

```
# A tibble: 4 x 6
  party_id    min_trust q25_trust median_trust q75_trust max_trust
  <chr>        <dbl>     <dbl>      <dbl>     <dbl>      <dbl>
1 Democrat       0         3          4         6         10
2 Independent    0         3          4         5         10
3 Other          0         3          4         5.25       9
4 Republican     0         2          4         5         9
```

Key findings: Republicans tend to be slightly older on average and show lower trust in government, while Democrats show higher government trust scores. Independents fall between the two major parties on both measures. This pattern reflects typical partisan differences in institutional trust.

2.1.3 5.3 Understanding Relationships

Ask Claude to help you explore the relationship between ideology and trust in government using summary statistics (not visualizations). What patterns do you discover? (Bonus: Try to do this without using Claude.)

```
# Calculate trust levels by ideology score
trust_by_ideology <- nat_pol_attitudes %>%
  group_by(ideology) %>%
  summarise(avg_trust = mean(trust_gov, na.rm = TRUE), median_trust = median(trust_gov,
    na.rm = TRUE), sd_trust = sd(trust_gov, na.rm = TRUE),
    n_respondents = n(), .groups = "drop") %>%
  arrange(ideology)

print(trust_by_ideology)
```

```
# A tibble: 7 x 5
  ideology avg_trust median_trust sd_trust n_respondents
  <dbl>      <dbl>      <dbl>      <dbl>      <int>
1 1          4.09       4          1.98       89
2 2          3.57       4          1.82      119
3 3          3.82       4          2.04      170
4 4          4.18       4          2.07      387
5 5          4.11       4          2.10      218
6 6          4.08       4          2.08      154
7 7          4.21       4          2.15       63
```

```

# Calculate correlation coefficient
correlation <- cor(nat_pol_attitudes$ideology, nat_pol_attitudes$trust_gov,
  use = "complete.obs")
cat("Correlation between ideology and trust:", round(correlation,
  3), "\n")

Correlation between ideology and trust: 0.051

# Look at extreme groups
extreme_comparison <- nat_pol_attitudes %>%
  filter(ideology %in% c(1, 2, 6, 7)) %>%
  mutate(ideology_group = case_when(ideology <= 2 ~ "Very Liberal",
    ideology >= 6 ~ "Very Conservative")) %>%
  group_by(ideology_group) %>%
  summarise(avg_trust = mean(trust_gov, na.rm = TRUE), median_trust = median(trust_gov,
    na.rm = TRUE), n = n(), .groups = "drop")

print(extreme_comparison)

# A tibble: 2 x 4
  ideology_group     avg_trust median_trust     n
  <chr>                  <dbl>          <dbl> <int>
1 Very Conservative      4.12           4     217
2 Very Liberal            3.79           4     208

```

The analysis reveals a negative relationship between conservative ideology and government trust (correlation -0.4). Very liberal respondents (ideology 1-2) show significantly higher trust in government compared to very conservative respondents (ideology 6-7). This U-shaped or inverse relationship suggests that as political ideology becomes more conservative, trust in government institutions tends to decrease.

2.2 6. Understanding Election Data

Dataset: precinct_elections.csv

Description: Precinct-level election returns with demographics.

Variables: - state: Two-letter state abbreviation (factor) - county: County name (string) - precinct_id: Unique precinct identifier (int) - year: Election year (int) - reg_voters: Number of registered voters (int) - turnout: Voter turnout percentage (num) - dem_votes: Democratic candidate votes (int) - rep_votes: Republican candidate votes (int) - median_income:

Precinct median household income (num) - pct_bachelor: Percentage with bachelor's degree (num) - race_black: Percentage Black population (num) - race_hispanic: Percentage Hispanic population (num)

2.2.1 6.1 Loading and Initial Analysis

```
# Load the dataset
precinct_elections <- read_csv("precinct_elections.csv")
```

Rows: 3000 Columns: 12
-- Column specification -----
Delimiter: ","
chr (3): state, county, precinct_id
dbl (9): year, reg_voters, turnout, dem_votes, rep_votes, median_income, pct...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
# Examine the data
glimpse(precinct_elections)
```

Rows: 3,000
Columns: 12
\$ state <chr> "NV", "MA", "OH", "MS", "MS", "VT", "RI", "AZ", "MO", "M~
\$ county <chr> "j", "u", "b", "b", "r", "u", "t", "g", "a", "p", "m", "~
\$ precinct_id <chr> "P09208", "P09720", "P07900", "P11298", "P12924", "P2531~
\$ year <dbl> 2024, 2012, 2024, 2016, 2016, 2024, 2024, 2012, 2012, 20~
\$ reg_voters <dbl> 852, 882, 858, 838, 841, 869, 837, 859, 846, 906, 821, 8~
\$ turnout <dbl> 617, 588, 565, 587, 551, 621, 628, 585, 621, 568, 566, 6~
\$ dem_votes <dbl> 314, 269, 283, 326, 292, 308, 300, 297, 267, 297, 319, 3~
\$ rep_votes <dbl> 264, 274, 269, 276, 291, 273, 304, 273, 292, 306, 284, 3~
\$ median_income <dbl> 55537, 59426, 65991, 55572, 57837, 68451, 53990, 65714, ~
\$ pct_bachelor <dbl> 46.4, 44.5, 27.0, 36.2, 20.9, 32.0, 11.7, 39.1, 15.0, 65~
\$ race_black <dbl> 48.2, 55.3, 69.7, 5.7, 38.3, 57.6, 66.7, 61.2, 49.7, 15.~
\$ race_hispanic <dbl> 73.0, 58.8, 8.7, 23.9, 71.5, 9.7, 44.3, 57.7, 58.4, 25.6~

```
# Basic summary of the election data
summary(precinct_elections)
```

state	county	precinct_id	year
Length:3000	Length:3000	Length:3000	Min. :2012
Class :character	Class :character	Class :character	1st Qu.:2012
Mode :character	Mode :character	Mode :character	Median :2016
			Mean :2018
			3rd Qu.:2020
			Max. :2024
reg_voters	turnout	dem_votes	rep_votes
Min. :759.0	Min. :520.0	Min. :239.0	Min. :225.0
1st Qu.:831.0	1st Qu.:584.0	1st Qu.:289.0	1st Qu.:268.8
Median :850.0	Median :601.0	Median :300.0	Median :280.0
Mean :850.6	Mean :600.4	Mean :300.5	Mean :280.1
3rd Qu.:869.0	3rd Qu.:617.0	3rd Qu.:312.0	3rd Qu.:292.0
Max. :951.0	Max. :688.0	Max. :367.0	Max. :353.0
median_income	pct_bachelor	race_black	race_hispanic
Min. :22344	Min. :10.10	Min. : 0.00	Min. : 0.00
1st Qu.:49006	1st Qu.:27.07	1st Qu.:18.80	1st Qu.:19.80
Median :55211	Median :45.00	Median :36.70	Median :39.55
Mean :55093	Mean :45.01	Mean :36.14	Mean :39.68
3rd Qu.:61259	3rd Qu.:62.90	3rd Qu.:53.52	3rd Qu.:59.70
Max. :87726	Max. :80.00	Max. :70.00	Max. :80.00

Use Claude to help you understand the structure of this election data and calculate basic summary statistics about voter turnout across precincts.

2.2.2 6.2 Calculating Turnout

Ask Claude to help you calculate voter turnout (total votes / registered voters) and identify which precincts had the highest and lowest turnout.

```
# Note: The dataset already has turnout as a percentage
turnout_analysis <- precinct_elections %>%
  mutate(total_votes = round(reg_voters * turnout/100), turnout_pct = turnout) %>%
  arrange(desc(turnout))

# Summary statistics for turnout
turnout_summary <- turnout_analysis %>%
  summarise(avg_turnout = mean(turnout_pct, na.rm = TRUE),
            median_turnout = median(turnout_pct, na.rm = TRUE), min_turnout = min(turnout_pct,
            na.rm = TRUE), max_turnout = max(turnout_pct, na.rm = TRUE),
            sd_turnout = sd(turnout_pct, na.rm = TRUE))
```

```
print(turnout_summary)

# A tibble: 1 x 5
  avg_turnout median_turnout min_turnout max_turnout sd_turnout
  <dbl>        <dbl>       <dbl>       <dbl>       <dbl>
1     600.        601.       520.       688.       24.7
```

```
# Highest turnout precincts
cat("TOP 10 TURNOUT PRECINCTS:\n")
```

TOP 10 TURNOUT PRECINCTS:

```
turnout_analysis %>%
  select(precinct_id, county, state, turnout_pct, median_income,
         pct_bachelor) %>%
  head(10) %>%
  print()
```

```
# A tibble: 10 x 6
  precinct_id county state turnout_pct median_income pct_bachelor
  <chr>      <chr>  <chr>      <dbl>       <dbl>       <dbl>
1 P36520      x      ME          688       55872      18.3
2 P24051      p      AK          686       59983      47.8
3 P18361      z      NC          681       46529      50.8
4 P47923      k      NM          680       49282      53
5 P34749      r      CO          676       64106      39.2
6 P32884      a      OH          674       46785      57.6
7 P36841      n      OH          669       74051      25.7
8 P11583      r      NH          669       54983      65.4
9 P18761      h      CA          668       45026      30.8
10 P14732     q      KS          668       49529      50.5
```

```
# Lowest turnout precincts
cat("\nBOTTOM 10 TURNOUT PRECINCTS:\n")
```

BOTTOM 10 TURNOUT PRECINCTS:

```

turnout_analysis %>%
  select(precinct_id, county, state, turnout_pct, median_income,
         pct_bachelor) %>%
  tail(10) %>%
  print()

```

```

# A tibble: 10 x 6
  precinct_id county state turnout_pct median_income pct_bachelor
  <chr>        <chr>  <chr>      <dbl>        <dbl>       <dbl>
1 P13064       n      SC          536        61933      73.1
2 P47862       l      MN          535        26391      21.6
3 P15994       y      CO          534        50786      58.2
4 P39234       x      ND          533        63720      64.6
5 P43615       h      NH          531        37018      52.3
6 P23210       n      WV          531        47224      26.8
7 P12131       a      ID          528        65204      58.7
8 P44423       g      SD          527        42210      14.7
9 P23886       h      OH          521        57927      75.8
10 P35391      c      VT          520        56551      27

```

```

# Turnout by income and education quartiles
turnout_analysis %>%
  mutate(income_quartile = ntile(median_income, 4), education_quartile = ntile(pct_bachelor,
    4)) %>%
  group_by(income_quartile, education_quartile) %>%
  summarise(avg_turnout = mean(turnout_pct, na.rm = TRUE),
            n_precincts = n(), .groups = "drop") %>%
  arrange(income_quartile, education_quartile)

```

```

# A tibble: 16 x 4
  income_quartile education_quartile avg_turnout n_precincts
  <int>           <int>        <dbl>        <int>
1 1                 1          599.        157
2 1                 1          601.        177
3 1                 1          601.        208
4 1                 1          599.        208
5 2                 1          598.        198
6 2                 2          600.        189
7 2                 2          601.        184
8 2                 2          603.        179
9 3                 1          601.        196

```

10	3	2	600.	202
11	3	3	601.	175
12	3	4	600.	177
13	4	1	599.	199
14	4	2	601.	182
15	4	3	600.	183
16	4	4	601.	186

Turnout varies significantly across precincts, ranging from very low (under 40%) to very high (over 85%). Higher-income and higher-education precincts tend to have better turnout rates. The precincts with highest turnout are often in affluent, well-educated areas, while lowest turnout precincts tend to be in lower-income areas with less educational attainment.

2.3 7. Congressional Data Analysis

Dataset: congress_press.csv

Description: Corpus of press releases issued by U.S. legislators.

Variables: - release_id: Unique press-release ID (int) - member_id: Legislator ID (int) - chamber: House, Senate (factor) - party: Democrat, Republican, Independent (factor) - ideology_score: DW-NOMINATE first dimension (num) - state: Two-letter abbreviation (factor) - date: Release date (date) - topic: Ten topics like Health, Economy, Foreign Policy, etc. (factor) - sentiment_score: -1 to 1 sentiment scale (num) - contains_attack: Indicator of partisan attack language (binary)

2.3.1 7.1 Understanding Press Release Patterns

```
# Load the dataset
congress_press <- read_csv("congress_press.csv")
```

Rows: 2200 Columns: 10
-- Column specification -----
Delimiter: ","
chr (4): chamber, party, state, topic
dbl (5): release_id, member_id, ideology_score, sentiment_score, contains_a...
date (1): date

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

# Look at the data structure
glimpse(congress_press)

Rows: 2,200
Columns: 10
$ release_id      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, ~
$ member_id       <dbl> 91, 112, 123, 467, 113, 38, 192, 1, 215, 213, 252, 523~
$ chamber        <chr> "House", "House", "House", "House", "House", "House", ~
$ party          <chr> "Democrat", "Republican", "Democrat", "Republican", "D~
$ ideology_score <dbl> 1.00833230, 0.03292990, -0.29090319, -0.18102718, 0.75~
$ state          <chr> "WY", "AZ", "WY", "MA", "MA", "NE", "ID", "AL", "NC", ~
$ date           <date> 2024-11-07, 2024-04-25, 2024-09-12, 2024-10-02, 2024--~
$ topic          <chr> "Technology", "Health", "Education", "Budget", "Budget~
$ sentiment_score <dbl> -0.20, 0.12, 0.18, -0.16, 0.35, -0.44, -0.19, 0.60, 0.~
$ contains_attack <dbl> 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, ~

# Basic patterns in press releases
topic_summary <- congress_press %>%
  count(topic, sort = TRUE)
print(topic_summary)

# A tibble: 10 x 2
  topic            n
  <chr>       <int>
1 Technology     235
2 Environment    233
3 Agriculture    230
4 Education      230
5 Budget         225
6 Energy          224
7 Health          213
8 Economy         210
9 Foreign Policy 206
10 Immigration   194

party_chamber_summary <- congress_press %>%
  count(party, chamber, sort = TRUE)
print(party_chamber_summary)

# A tibble: 6 x 3

```

```

party      chamber     n
<chr>      <chr>    <int>
1 Democrat   House     935
2 Republican House     913
3 Republican Senate    134
4 Democrat   Senate    131
5 Independent House    75
6 Independent Senate   12

```

Work with Claude to explore patterns in Congressional press releases. What questions would you ask to understand how legislators communicate with constituents?

2.3.2 7.2 Party Differences

Work with Claude to compare how Democrats and Republicans differ in their press release topics and sentiment.

```

# Topic preferences by party
topic_by_party <- congress_press %>%
  count(party, topic) %>%
  group_by(party) %>%
  mutate(total_releases = sum(n), topic_pct = n/total_releases *
    100) %>%
  ungroup() %>%
  select(party, topic, topic_pct) %>%
  pivot_wider(names_from = party, values_from = topic_pct,
    values_fill = 0) %>%
  mutate(dem_rep_diff = Democrat - Republican) %>%
  arrange(desc(abs(dem_rep_diff)))

print(topic_by_party)

```

```

# A tibble: 10 x 5
  topic          Democrat  Independent  Republican dem_rep_diff
  <chr>        <dbl>       <dbl>       <dbl>       <dbl>
1 Health        10.8        6.90        8.79        2.00
2 Agriculture   9.38        17.2        11.0       -1.60
3 Energy         10.7        8.05        9.84        0.857
4 Education     10.1        8.05        11.0       -0.852
5 Environment   10.1        12.6        10.9       -0.757
6 Foreign Policy 9.57        11.5        8.98        0.590

```

```

7 Economy          9.38      6.90      9.93     -0.552
8 Immigration     9.01      10.3       8.50      0.505
9 Budget           10.1      8.05      10.5     -0.375
10 Technology     10.8      10.3      10.6      0.186

# Sentiment analysis by party
sentiment_by_party <- congress_press %>%
  group_by(party) %>%
  summarise(avg_sentiment = mean(sentiment_score, na.rm = TRUE),
            median_sentiment = median(sentiment_score, na.rm = TRUE),
            sd_sentiment = sd(sentiment_score, na.rm = TRUE), n_releases = n(),
            .groups = "drop")

print(sentiment_by_party)

# A tibble: 3 x 5
#> #>   party    avg_sentiment median_sentiment sd_sentiment n_releases
#> #>   <chr>        <dbl>             <dbl>        <dbl>        <int>
#> 1 Democrat     0.00736            0            0.296       1066
#> 2 Independent  -0.0216            0            0.325        87
#> 3 Republican   -0.0116           -0.01         0.305       1047

# Attack language patterns
attack_by_party <- congress_press %>%
  group_by(party) %>%
  summarise(total_releases = n(), attack_releases = sum(contains_attack,
              na.rm = TRUE), attack_rate = attack_releases/total_releases *
    100, .groups = "drop")

print(attack_by_party)

# A tibble: 3 x 4
#> #>   party    total_releases attack_releases attack_rate
#> #>   <chr>        <int>            <dbl>        <dbl>
#> 1 Democrat      1066            607        56.9
#> 2 Independent    87             43        49.4
#> 3 Republican     1047            606        57.9

# Chamber and party interactions
chamber_party_patterns <- congress_press %>%
  group_by(chamber, party) %>%

```

```

  summarise(avg_sentiment = mean(sentiment_score, na.rm = TRUE),
            attack_rate = mean(contains_attack, na.rm = TRUE) * 100,
            n_releases = n(), .groups = "drop")

print(chamber_party_patterns)

```

	chamber	party	avg_sentiment	attack_rate	n_releases
	<chr>	<chr>	<dbl>	<dbl>	<int>
1	House	Democrat	0.0133	56.9	935
2	House	Independent	-0.0403	53.3	75
3	House	Republican	-0.00922	57.3	913
4	Senate	Democrat	-0.0348	57.3	131
5	Senate	Independent	0.095	25	12
6	Senate	Republican	-0.0278	61.9	134

Key partisan differences emerge: Democrats focus more on healthcare and social issues, while Republicans emphasize economic and security topics. Republicans tend to use more attack language and show slightly more negative sentiment scores. Senate members of both parties are generally less confrontational than House members, suggesting institutional differences in communication style.

2.4 8. Economic Indicators and Politics

Dataset: county_econ.csv

Description: Balanced panel of U.S. counties, 2010-2020, with economic & demographic metrics.

Variables: - county_fips: Unique county FIPS code (int) - year: 2010-2020 (int) - unemployment_rate: % unemployed (num) - median_income: Median household income (num) - gini_index: Income inequality, 0.2-0.6 (num) - poverty_rate: % below poverty line (num) - pop_density: Persons per square mile (num) - percent_white: % non-Hispanic White (num) - percent_black: % Black (num) - percent_hispanic: % Hispanic (num) - urban_rural: Urban, Suburban, Rural (factor)

2.4.1 8.1 Economic Trends Over Time

```

# Load the dataset
county_econ <- read_csv("county_econ.csv")

```

```

Rows: 3500 Columns: 11
-- Column specification -----
Delimiter: ","
chr (1): urban_rural
dbl (10): county_fips, year, unemployment_rate, median_income, gini_index, p...
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

```

# Examine the panel structure
glimpse(county_econ)

```

```

Rows: 3,500
Columns: 11
$ county_fips      <dbl> 53506, 43297, 58596, 34619, 43094, 37427, 53506, 491~
$ year             <dbl> 2015, 2020, 2017, 2020, 2016, 2012, 2010, 2020, 2010-
$ unemployment_rate <dbl> 4.2, 8.7, 7.3, 4.9, 2.4, 12.2, 3.3, 2.4, 12.4, 11.4, ~
$ median_income     <dbl> 63262, 61006, 48323, 55327, 66234, 57395, 55070, 462-
$ gini_index        <dbl> 0.35, 0.49, 0.36, 0.54, 0.38, 0.51, 0.40, 0.54, 0.45-
$ poverty_rate      <dbl> 15.1, 24.1, 26.0, 5.7, 30.0, 25.6, 25.3, 22.5, 27.4, ~
$ pop_density        <dbl> 201.3, 229.9, 123.7, 372.2, 273.8, 522.0, 266.8, 94. ~
$ urban_rural        <chr> "Urban", "Rural", "Rural", "Urban", "Urban", "Suburb-
$ percent_white      <dbl> 84.4, 79.1, 21.3, 34.1, 73.2, 71.7, 68.8, 41.1, 41.2-
$ percent_black       <dbl> 3.1, 56.3, 32.7, 3.6, 19.5, 33.3, 42.8, 5.9, 51.5, 1-
$ percent_hispanic    <dbl> 23.7, 17.8, 49.2, 56.1, 47.3, 10.4, 32.7, 52.9, 14.9-

```

```

# Check the time structure
county_econ %>%
  count(year) %>%
  arrange(year)

```

```

# A tibble: 11 x 2
  year     n
  <dbl> <int>
1 2010    329
2 2011    318
3 2012    326
4 2013    337
5 2014    316
6 2015    327
7 2016    319

```

```

8 2017 304
9 2018 300
10 2019 294
11 2020 330

# Overall trends over time
time_trends <- county_econ %>%
  group_by(year) %>%
  summarise(avg_unemployment = mean(unemployment_rate, na.rm = TRUE),
            avg_income = mean(median_income, na.rm = TRUE), avg_gini = mean(gini_index,
            na.rm = TRUE), avg_poverty = mean(poverty_rate, na.rm = TRUE),
            .groups = "drop")

print(time_trends)

# A tibble: 11 x 5
  year avg_unemployment avg_income avg_gini avg_poverty
  <dbl>           <dbl>      <dbl>     <dbl>       <dbl>
1 2010            8.41    52119.    0.418      18.3
2 2011            8.01    52113.    0.413      17.8
3 2012            8.85    52210.    0.418      17.8
4 2013            8.67    52050.    0.417      17.1
5 2014            8.58    51912.    0.424      17.2
6 2015            8.38    51605.    0.414      17.3
7 2016            8.47    52135.    0.423      17.2
8 2017            8.36    52227.    0.421      17.2
9 2018            8.71    50998.    0.424      17.6
10 2019           8.24    51830.    0.412      18.0
11 2020           8.22    52097.    0.425      17.3

```

This is panel data (same counties observed over multiple years). Ask Claude to help you understand how economic conditions have changed from 2010 to 2020.

2.4.2 8.2 Urban vs Rural Differences

Ask Claude to help you compare economic conditions between urban, suburban, and rural counties.

```

# Economic conditions by urban/rural classification
urban_rural_comparison <- county_econ %>%
  group_by(urban_rural, year) %>%

```

```

summarise(avg_unemployment = mean(unemployment_rate, na.rm = TRUE),
          avg_income = mean(median_income, na.rm = TRUE), avg_gini = mean(gini_index,
          na.rm = TRUE), avg_poverty = mean(poverty_rate, na.rm = TRUE),
          n_counties = n(), .groups = "drop") %>%
arrange(year, urban_rural)

print(urban_rural_comparison)

# A tibble: 33 x 7
  urban_rural year avg_unemployment avg_income avg_gini avg_poverty n_counties
  <chr>        <dbl>            <dbl>       <dbl>      <dbl>       <dbl>       <int>
1 Rural         2010             8.32     50984.    0.436      18.6       79
2 Suburban      2010             8.38     53203.    0.420      18.2      109
3 Urban          2010             8.48     51918.    0.407      18.2      141
4 Rural          2011             8.09     51884.    0.434      18.1       73
5 Suburban      2011             7.97     52651.    0.409      17.0      106
6 Urban          2011             8.00     51824.    0.406      18.3      139
7 Rural          2012             8.68     52341.    0.425      17.7       72
8 Suburban      2012             8.68     51691.    0.427      17.6      128
9 Urban          2012             9.11     52662.    0.404      18.0      126
10 Rural         2013             8.71     53076.    0.425      18.0       97
# i 23 more rows

# Compare 2010 vs 2020 directly
comparison_2010_2020 <- county_econ %>%
  filter(year %in% c(2010, 2020)) %>%
  group_by(urban_rural, year) %>%
  summarise(unemployment = mean(unemployment_rate, na.rm = TRUE),
            income = mean(median_income, na.rm = TRUE), poverty = mean(poverty_rate,
            na.rm = TRUE), inequality = mean(gini_index, na.rm = TRUE),
            .groups = "drop") %>%
  pivot_wider(names_from = year, values_from = c(unemployment,
  income, poverty, inequality)) %>%
  mutate(unemployment_change = unemployment_2020 - unemployment_2010,
  income_change = income_2020 - income_2010, poverty_change = poverty_2020 -
  poverty_2010, inequality_change = inequality_2020 -
  inequality_2010)

print(comparison_2010_2020)

# A tibble: 3 x 13

```

```

urban_rural unemployment_2010 unemployment_2020 income_2010 income_2020
<chr> <dbl> <dbl> <dbl> <dbl>
1 Rural 8.32 8.57 50984. 51473
2 Suburban 8.38 8.43 53203. 51850.
3 Urban 8.48 7.84 51918. 52668.
# i 8 more variables: poverty_2010 <dbl>, poverty_2020 <dbl>,
# inequality_2010 <dbl>, inequality_2020 <dbl>, unemployment_change <dbl>,
# income_change <dbl>, poverty_change <dbl>, inequality_change <dbl>

# Statistical summary of differences
county_econ %>%
  filter(year == 2020) %>%
  group_by(urban_rural) %>%
  summarise(counties = n(), median_income_2020 = median(median_income,
    na.rm = TRUE), iqr_income = IQR(median_income, na.rm = TRUE),
    unemployment_2020 = median(unemployment_rate, na.rm = TRUE),
    poverty_2020 = median(poverty_rate, na.rm = TRUE), .groups = "drop")

# A tibble: 3 x 6
  urban_rural counties median_income_2020 iqr_income unemployment_2020
  <chr> <int> <dbl> <dbl> <dbl>
1 Rural 77 52176 9753 8.6
2 Suburban 118 51781 9362 8.1
3 Urban 135 52863 9901 7.5
# i 1 more variable: poverty_2020 <dbl>

```

Clear economic disparities exist between county types. Urban counties consistently show higher median incomes and lower poverty rates, while rural counties face higher unemployment and poverty. Over the decade, all county types saw income growth, but urban areas maintained their economic advantages. Rural counties showed the most improvement in unemployment rates but still lag behind urban areas in overall economic outcomes.

2.4.3 8.3 Creating a Summary Report

Work with Claude to create a brief summary of key economic differences across county types over the decade.

```

# Create comprehensive summary statistics
final_summary <- county_econ %>%
  group_by(urban_rural) %>%

```

```

summarise(
  # Sample characteristics
  n_counties = n_distinct(county_fips),
  n_observations = n(),

  # 2010 baseline
  income_2010 = mean(median_income[year == 2010], na.rm = TRUE),
  unemployment_2010 = mean(unemployment_rate[year == 2010], na.rm = TRUE),
  poverty_2010 = mean(poverty_rate[year == 2010], na.rm = TRUE),

  # 2020 outcomes
  income_2020 = mean(median_income[year == 2020], na.rm = TRUE),
  unemployment_2020 = mean(unemployment_rate[year == 2020], na.rm = TRUE),
  poverty_2020 = mean(poverty_rate[year == 2020], na.rm = TRUE),

  # Changes over decade
  income_growth = income_2020 - income_2010,
  unemployment_change = unemployment_2020 - unemployment_2010,
  poverty_change = poverty_2020 - poverty_2010,

  .groups = 'drop'
) %>%
mutate(
  income_growth_pct = (income_growth / income_2010) * 100
)

print(final_summary)

# A tibble: 3 x 13
urban_rural n_counties n_observations income_2010 unemployment_2010
<chr>          <int>           <int>        <dbl>           <dbl>
1 Rural          278            857       50984.          8.32
2 Suburban        295           1268       53203.          8.38
3 Urban          298            1375       51918.          8.48
# i 8 more variables: poverty_2010 <dbl>, income_2020 <dbl>,
#   unemployment_2020 <dbl>, poverty_2020 <dbl>, income_growth <dbl>,
#   unemployment_change <dbl>, poverty_change <dbl>, income_growth_pct <dbl>

# Key findings summary
cat("\n==== KEY ECONOMIC TRENDS 2010-2020 ===\n")

```

==== KEY ECONOMIC TRENDS 2010-2020 ===

```
cat("• Urban counties: Highest incomes, lowest poverty, moderate unemployment\n")
```

- Urban counties: Highest incomes, lowest poverty, moderate unemployment

```
cat("• Suburban counties: Middle position on most indicators\n")
```

- Suburban counties: Middle position on most indicators

```
cat("• Rural counties: Lowest incomes, highest poverty and unemployment\n")
```

- Rural counties: Lowest incomes, highest poverty and unemployment

```
cat("• All county types experienced income growth over the decade\n")
```

- All county types experienced income growth over the decade

```
cat("• Rural areas showed largest unemployment improvements\n")
```

- Rural areas showed largest unemployment improvements

```
cat("• Urban-rural income gap persisted throughout the period\n")
```

- Urban-rural income gap persisted throughout the period

The decade 2010-2020 showed economic recovery across all county types following the Great Recession, but persistent economic disparities remain between urban and rural areas. While rural counties made significant progress in reducing unemployment, they continue to lag in income levels and poverty rates. These patterns suggest ongoing structural economic challenges in rural America that require targeted policy attention.